

Part B: Customer Churn Prediction

Analyze the provided customer data ("Customer_data") to build a machine learning model that predicts customer churn. The analysis should include data exploration, preprocessing (handling missing values, encoding, scaling), model training, and evaluation using metrics like accuracy, precision, recall, and F1 score. The goal is to identify customers at high risk of churning to inform retention strategies.

Load the data

Subtask:

Load the customer data from the specified file into a pandas DataFrame.

Reasoning: Import pandas and load the data into a DataFrame.

```
import pandas as pd
```

```
df = pd.read_csv('Customer_data.csv')
df.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...

5 rows × 21 columns

Explore and preprocess the data

Subtask:

Perform exploratory data analysis, handle missing values, encode categorical features, and scale numerical features to prepare the data for modeling.

Reasoning: Examine the data types and summary statistics, identify and handle missing values, identify and encode categorical features, and identify and scale numerical features to prepare the data for modeling.

```
# 1. Examine data types and summary statistics
display(df.info())
display(df.describe(include='all'))

# 2. Identify and handle missing values
# Check for missing values
display(df.isnull().sum())

# The 'TotalCharges' column has missing values.
# Since it's a numerical column, we can impute with the mean or median,
# or drop rows with missing values. Given the small number of missing values
# compared to the dataset size, dropping the rows is a reasonable approach.
df.dropna(inplace=True)

# Verify that missing values have been handled
display(df.isnull().sum())

# Convert 'TotalCharges' to numeric, coercing errors
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')

# Drop rows where 'TotalCharges' became NaN after coercion (if any)
df.dropna(inplace=True)
```

```
# 3. Identify categorical features and apply one-hot encoding
categorical_features = df.select_dtypes(include='object').columns
# Exclude 'customerID' as it is an identifier
categorical_features = categorical_features.drop('customerID')

df = pd.get_dummies(df, columns=categorical_features, drop_first=True)

# 4. Identify numerical features and scale them
# Numerical features are now those that are not 'customerID' and are not boolean (from one-hot encoding)
numerical_features = df.select_dtypes(include=['int64', 'float64']).columns

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df[numerical_features] = scaler.fit_transform(df[numerical_features])

display(df.head())
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
Column Non-Null Count Dtype

0 customerID 7043 non-null object
1 gender 7043 non-null object
2 SeniorCitizen 7043 non-null int64
3 Partner 7043 non-null object
4 Dependents 7043 non-null object
5 tenure 7043 non-null int64
6 PhoneService 7043 non-null object
7 MultipleLines 7043 non-null object
8 InternetService 7043 non-null object
9 OnlineSecurity 7043 non-null object
10 OnlineBackup 7043 non-null object
11 DeviceProtection 7043 non-null object
12 TechSupport 7043 non-null object
13 StreamingTV 7043 non-null object
14 StreamingMovies 7043 non-null object
15 Contract 7043 non-null object
16 PaperlessBilling 7043 non-null object
17 PaymentMethod 7043 non-null object
18 MonthlyCharges 7043 non-null float64
19 TotalCharges 7032 non-null float64
20 Churn 7043 non-null object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
None

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineS
count	7043	7043	7043.000000	7043	7043	7043.000000	7043	7043	7043	
unique	7043	2	NaN	2	2	NaN	2	3	3	
top	3186-AJIEK	Male	NaN	No	No	NaN	Yes	No	Fiber optic	
freq	1	3555	NaN	3641	4933	NaN	6361	3390	3096	
mean	NaN	NaN	0.162147	NaN	NaN	32.371149	NaN	NaN	NaN	
std	NaN	NaN	0.368612	NaN	NaN	24.559481	NaN	NaN	NaN	
min	NaN	NaN	0.000000	NaN	NaN	0.000000	NaN	NaN	NaN	
25%	NaN	NaN	0.000000	NaN	NaN	9.000000	NaN	NaN	NaN	
50%	NaN	NaN	0.000000	NaN	NaN	29.000000	NaN	NaN	NaN	
75%	NaN	NaN	0.000000	NaN	NaN	55.000000	NaN	NaN	NaN	
max	NaN	NaN	1.000000	NaN	NaN	72.000000	NaN	NaN	NaN	

11 rows × 21 columns

0

customerID

gender

SeniorCitizen

Partner

Dependents

tenure

PhoneService

MultipleLines

InternetService

OnlineSecurity

OnlineBackup

DeviceProtection

TechSupport

StreamingTV

StreamingMovies

Contract

PaperlessBilling

PaymentMethod

MonthlyCharges

```

TotalCharges    11
Churn           0

dtype: int64
0
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0

```

```
dtype: int64
```

	customerID	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender_Male	Partner_Yes	Dependents_Yes	PhoneService_Yes
0	7590-VHVEG	-0.440327	-1.280248	-1.161694	-0.994194	False	True	False	False
1	5575-GNVDE	-0.440327	0.064303	-0.260878	-0.173740	True	False	False	True
2	3668-QPYBK	-0.440327	-1.239504	-0.363923	-0.959649	True	False	False	True
3	7795-CFOCW	-0.440327	0.512486	-0.747850	-0.195248	True	False	False	False
4	9237-HQITU	-0.440327	-1.239504	0.196178	-0.940457	False	False	False	True

```
5 rows x 32 columns
```

✓ Split the data

Subtask:

Split the data into training and testing sets.

Reasoning: Split the data into training and testing sets using `train_test_split`.

```
from sklearn.model_selection import train_test_split

X = df.drop(['Churn_Yes', 'customerID'], axis=1)
y = df['Churn_Yes']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of y_test:", y_test.shape)
```

```
↗ Shape of X_train: (5625, 30)
  Shape of X_test: (1407, 30)
  Shape of y_train: (5625,)
  Shape of y_test: (1407,)
```

✓ Build and train a machine learning model

Subtask:

Choose a suitable classification model (e.g., Logistic Regression, Random Forest, or Gradient Boosting) and train it on the training data.

Reasoning: Import the `LogisticRegression` model, instantiate it, and train it using the training data.

```
from sklearn.linear_model import LogisticRegression

# Instantiate the Logistic Regression model
model = LogisticRegression(random_state=42)

# Train the model on the training data
model.fit(X_train, y_train)
```

```
↗ LogisticRegression ⓘ ?
  LogisticRegression(random_state=42)
```

✓ Evaluate the model

Subtask:

Evaluate the trained model's performance on the testing data using appropriate metrics such as accuracy, precision, recall, and F1 score.

Reasoning: Evaluate the trained model's performance on the testing data using appropriate metrics such as accuracy, precision, recall, and F1 score and print the results.

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Use the trained model to predict churn on the test data
y_pred = model.predict(X_test)

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Print the evaluation metrics
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")
```

```

Accuracy: 0.7875
Precision: 0.6206
Recall: 0.5160
F1 Score: 0.5635

```

Summary:

Data Analysis Key Findings

- The dataset initially contained 7043 entries and 21 columns.
- The 'TotalCharges' column had 11 missing values, which were handled by dropping the corresponding rows.
- Categorical features were successfully one-hot encoded, and numerical features ('SeniorCitizen', 'tenure', 'MonthlyCharges', and 'TotalCharges') were scaled using `StandardScaler`.
- The data was split into training (80%) and testing (20%) sets, resulting in 5625 training samples and 1407 testing samples.
- A Logistic Regression model was trained on the preprocessed training data.
- The trained model's performance on the test set was evaluated with the following metrics:
 - Accuracy: 78.75%
 - Precision: 62.06%
 - Recall: 51.60%
 - F1 Score: 56.35%

Insights or Next Steps

- The model demonstrates a reasonable overall accuracy, but the relatively lower recall suggests that it misses a significant portion of actual churn cases. For churn prediction, where identifying all potential churners is crucial for intervention, improving recall should be a focus.
- Further steps could involve experimenting with different classification models (e.g., Random Forest, Gradient Boosting), optimizing hyperparameters, or exploring techniques to handle the class imbalance if it exists in the churn data to potentially improve the recall and overall performance.

✓ Visualize the data

Subtask:

Create visualizations to explore the data and churn patterns.

Reasoning: Visualize the distribution of the target variable and some key features, and the relationship between features and the target variable to gain insights about the data and churn patterns.

```

import matplotlib.pyplot as plt
import seaborn as sns

# Set the style for the plots
sns.set_style("whitegrid")

# Visualize the distribution of the target variable 'Churn_Yes'
plt.figure(figsize=(6, 4))
sns.countplot(x='Churn_Yes', data=df)
plt.title('Distribution of Churn')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.show()

# Visualize the distribution of 'tenure'
plt.figure(figsize=(8, 5))
sns.histplot(data=df, x='tenure', kde=True)
plt.title('Distribution of Tenure')
plt.xlabel('Tenure (months)')
plt.ylabel('Frequency')
plt.show()

# Visualize the distribution of 'MonthlyCharges'
plt.figure(figsize=(8, 5))
sns.histplot(data=df, x='MonthlyCharges', kde=True)
plt.title('Distribution of Monthly Charges')
plt.xlabel('Monthly Charges')
plt.ylabel('Frequency')
plt.show()

# Visualize the distribution of 'TotalCharges'
plt.figure(figsize=(8, 5))

```

```
plt.figure(figsize=(10, 5))
sns.histplot(data=df, x='TotalCharges', kde=True)
plt.title('Distribution of Total Charges')
plt.xlabel('Total Charges')
plt.ylabel('Frequency')
plt.show()

# Visualize the relationship between 'tenure' and 'Churn_Yes'
plt.figure(figsize=(8, 5))
sns.boxplot(x='Churn_Yes', y='tenure', data=df)
plt.title('Churn vs. Tenure')
plt.xlabel('Churn')
plt.ylabel('Tenure (months)')
plt.show()

# Visualize the relationship between 'MonthlyCharges' and 'Churn_Yes'
plt.figure(figsize=(8, 5))
sns.boxplot(x='Churn_Yes', y='MonthlyCharges', data=df)
plt.title('Churn vs. Monthly Charges')
plt.xlabel('Churn')
plt.ylabel('Monthly Charges')
plt.show()

# Visualize the relationship between 'TotalCharges' and 'Churn_Yes'
plt.figure(figsize=(8, 5))
sns.boxplot(x='Churn_Yes', y='TotalCharges', data=df)
plt.title('Churn vs. Total Charges')
plt.xlabel('Churn')
plt.ylabel('Total Charges')
plt.show()

# Visualize the relationship between 'Contract' and 'Churn_Yes'
plt.figure(figsize=(8, 5))
sns.countplot(x='Contract_Two year', hue='Churn_Yes', data=df)
plt.title('Churn vs. Contract Type')
plt.xlabel('Contract Type (Two Year)')
plt.ylabel('Count')
plt.xticks(ticks=[0, 1], labels=['Month-to-month/One Year', 'Two Year'])
plt.show()

# Visualize the relationship between 'InternetService' and 'Churn_Yes'
plt.figure(figsize=(8, 5))
sns.countplot(x='InternetService_Fiber optic', hue='Churn_Yes', data=df)
plt.title('Churn vs. Internet Service Type')
plt.xlabel('Internet Service Type (Fiber Optic)')
plt.ylabel('Count')
plt.xticks(ticks=[0, 1], labels=['DSL/No internet service', 'Fiber Optic'])
plt.show()
```



